


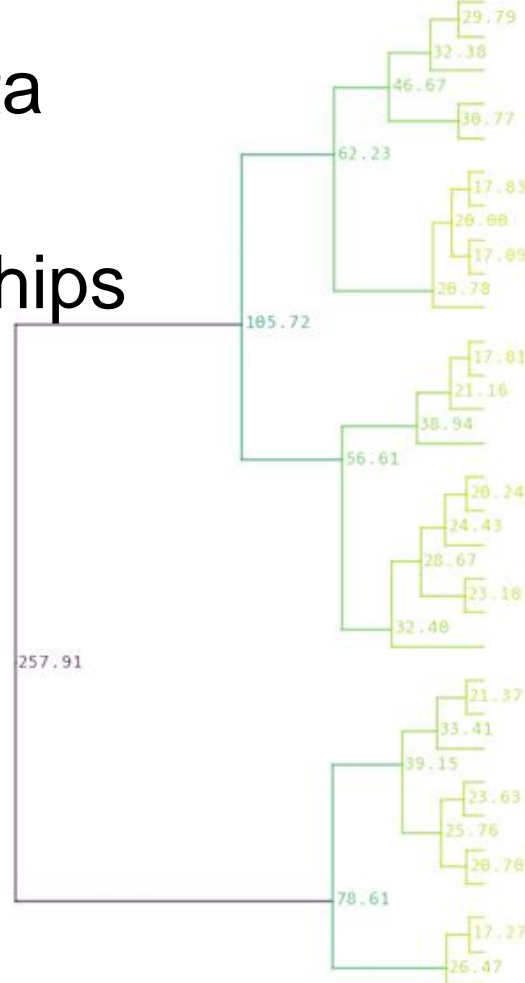
HIERARCHICAL CLUSTERING

Jesse Davis

Hierarchical Clustering

2

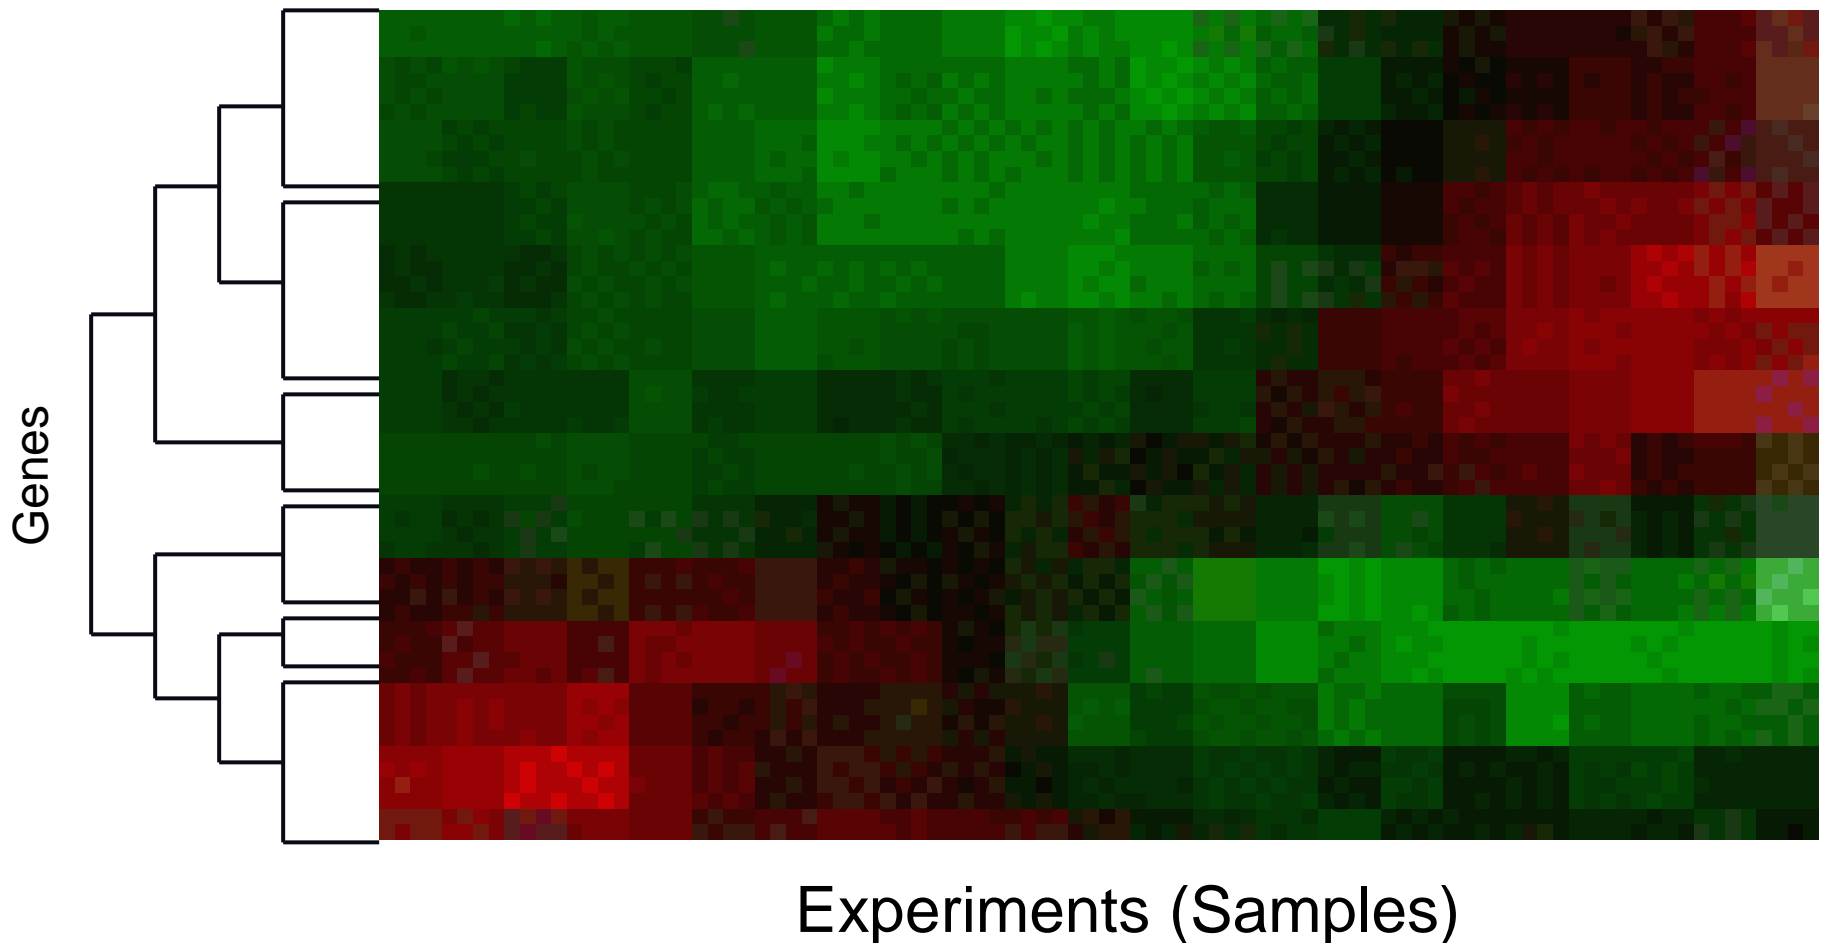
- Model hierarchical structure in the data
- Captures relationships among clusters



Example: Gene Expression

3

(Green = up-regulated, Red = down-regulated)




4

Basics

Hierarchical Clustering: Dendrogram

5



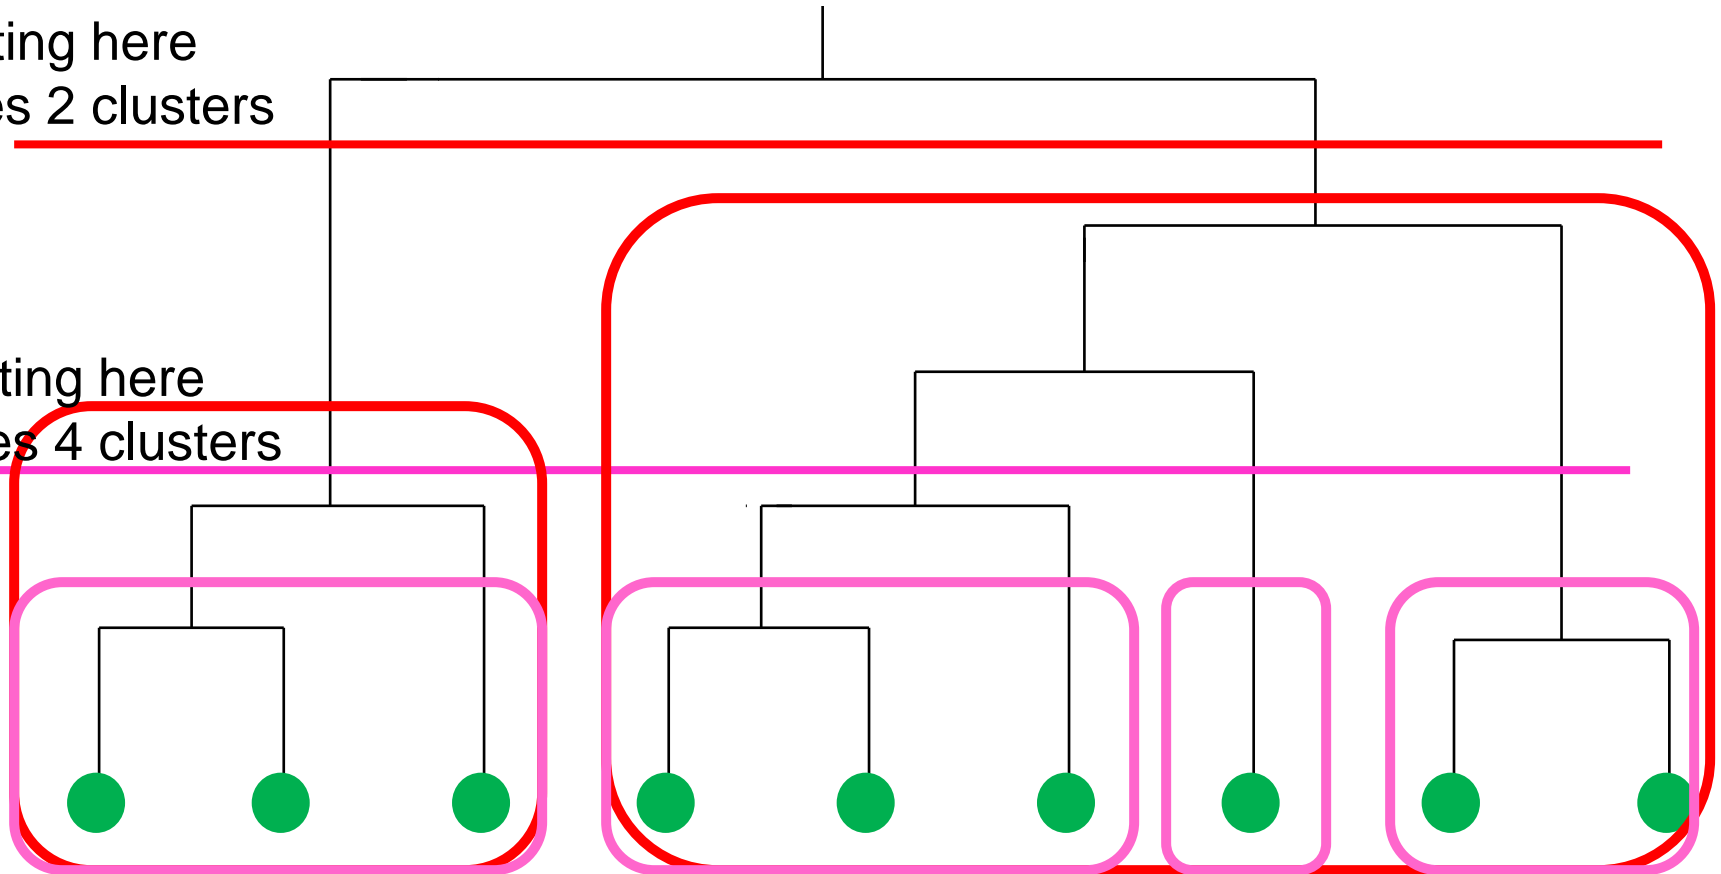
Note: Partitional Clustering from a Hierarchical Clustering

6

Can generate a partitional clustering from a hierarchical clustering by “cutting” the tree at some level

Cutting here
gives 2 clusters


Cutting here
gives 4 clusters




Hierarchical Clustering Approaches

7

- Top-down or divisive




- Bottom-up or agglomerative



Bottom-Up Example

8



Bottom-Up Hierarchical Clustering

9

Given: instances x_1, \dots, x_n

For $i = 1$ to n , $c_i = \{x_i\}$

$C = \{c_1, \dots, c_n\}$

$j = n$

While $|C| > 1$

$j = j+1$

$(c_a, c_b) = \operatorname{argmin} \operatorname{dist}(c_a, c_b)$

$c_j = c_a \cup c_b$

 add node to tree joining a and b

$C = (C - \{c_a, c_b\}) \cup c_j$

Return tree with root node j

Key question: Measuring distance

Distance Matrix

10

$D =$

0				
$d(2,1)$	0			
$d(3,1)$	$d(3,2)$	0		
\vdots	\vdots	\vdots	\ddots	
$d(n,1)$	$d(n,2)$	$d(n,3)$	0

Initial Distance Matrix for a Data Set with 5 Examples

11

- 1) Form five clusters, one for each example
- 2) Compute pairwise distance between initial clusters
(=pairwise distance between examples)

	1	2	3	4	5
1	0				
2	7.4	0			
3	0.7	7.1	0		
4	7.3	0.3	7.0	0	
5	0.5	6.9	0.6	6.8	0

Find Two Closest Clusters

12

	1	2	3	4	5
1	0				
2	7.4	0			
3	0.7	7.1	0		
4	7.3	0.3	7.0	0	
5	0.5	6.9	0.6	6.8	0

Smallest value in matrix

Update Distance Matrix

13


	1	(2,4)	3	5
1	0			
(2,4)	?	0		
3	0.7	?	0	
5	0.5	?	0.6	0

Question: What is the distance between the new cluster (2,4) and the other three clusters?


Measuring the Distance Between Two Clusters

14


- **Single link:** Distance of two most similar instances:
 $\text{dist}(c_u, c_v) = \min\{\text{dist}(a, b) \mid a \in c_u, b \in c_v\}$



- **Complete link:** Distance of two least similar instances:
 $\text{dist}(c_u, c_v) = \max\{\text{dist}(a, b) \mid a \in c_u, b \in c_v\}$



- **Average link:** Average distance between instances:
 $\text{dist}(c_u, c_v) = \text{avg}\{\text{dist}(a, b) \mid a \in c_u, b \in c_v\}$



Efficient Distance Updates

15

- If we merged c_u and c_v into c_j , we can determine distance to every other cluster:

- ▣ Single link:

$$\text{dist}(c_j, c_k) = \min(\text{dist}(c_u, c_k), \text{dist}(c_v, c_k))$$

- ▣ Complete link:

$$\text{dist}(c_j, c_k) = \max(\text{dist}(c_u, c_k), \text{dist}(c_v, c_k))$$

- ▣ Average link:

$$\text{dist}(c_j, c_k) = \frac{|c_u| * \text{dist}(c_u, c_k) + |c_v| * \text{dist}(c_v, c_k)}{|c_u| + |c_v|}$$

Note: The linkage choice is a hyper parameter for the bottom-up clustering algorithm

Illustrative Example Updates for Each Linkage Criteria

16

Single

	2	4	(2,4)
1	7.4	7.3	7.3
3	7.1	7.0	7.0
5	6.9	6.8	6.8

min

Complete

	2	4	(2,4)
1	7.4	7.3	7.4
3	7.1	7.0	7.1
5	6.9	6.8	6.9

max

Average


	2	4	(2,4)
1	7.4	7.3	7.35
3	7.1	7.0	7.05
5	6.9	6.8	6.85

*(Weighted)
mean*

Complete Link Dendrogram for Sample Dataset


17

$$\text{height} = h(c_j, c_k) = \frac{\text{dist}(c_j, c_k)}{2}$$




Single Link: Chaining

18




Single Link: Chaining

19



Single Link: Chaining


20



Single Link

21

□ Chaining:



□ Bottom line:

- ▣ Simple, fast
- ▣ Often low quality

Complete Link Hierarchical Summary

22

- Complexity: $O(n^3)$
 - ▣ $O(n^2)$ to build initial similarity matrix
 - ▣ $O(n)$ for the merges
- Fast algorithm: Requires $O(n^2)$ space
- Bottom line
 - ▣ Typically much faster than $O(n^3)$
 - ▣ Often good quality
 - ▣ No Chaining

23

Advanced Hierarchical Clustering

Other Hierarchical Clustering Methods

24

- Weaknesses of agglomerative clustering methods
 - ▣ **Do not scale well:** time complexity of at least $O(n^2)$, where n is the number of total objects
 - ▣ Can never undo what was done previously
- Integration of hierarchical with distance-based clustering
 - ▣ BIRCH: uses CF-tree and incrementally adjusts the quality of sub-clusters
 - ▣ CURE: selects well-scattered points from the cluster and then shrinks them towards the center of the cluster by a specified fraction

BIRCH: Balanced Iterative Reducing and Clustering using Hierarchies

25

- Incrementally construct a Clustering Feature (CF) tree
 - ▣ **Phase 1:** Scan DB to build an initial in-memory CF tree (each node: #points, sum, sum of squares)
 - ▣ **Phase 2:** use an arbitrary clustering algorithm to cluster the leaf nodes of the CF-tree
- ***Scales linearly:*** Finds a good clustering with a single scan
- ***Weaknesses:*** handles only numeric data, sensitive to order of data records

Definitions

26

□ **Centroid:** $\mu = \frac{1}{N} \sum_{i=1}^N x_i$

Average along
each dimension

- **Radius:** Average distance from member points to cluster centroid

$$R = \sqrt{\frac{\sum_{i=1}^N \sum_{j=1}^d (x_{ij} - \mu_j)^2}{N}}$$

- Captures tightness of cluster around centroid
- Note: Math and verbal definition not 100% aligned, but these come directly from the paper

Cluster Feature Vector

27

- Given: X_1, \dots, X_n , data points in a cluster where each with d -dimensions
- We define $CF = (N, LS, SS)$, where
 - ▣ N : Number of data points
 - ▣ $LS_j: \sum_{i=1}^n x_{ij}$
 - ▣ $SS_j: \sum_{i=1}^n x_{ij}^2$
- Note: CFs are additive!
 - ▣ E.g., $CF_1 + CF_2 = (N_1 + N_2, LS_1 + LS_2, SS_1 + SS_2)$

Cluster Feature Example

28

$$CF = (5, (16, 30), (54, 190))$$


(3,4)

(2,6)

(4,5)

(4,7)

(3,8)



$$LS_x = 3 + 2 + 4 + 4 + 3 = 16$$

$$LS_y = 4 + 6 + 5 + 7 + 8 = 30$$

$$SS_x = 3^2 + 2^2 + 4^2 + 4^2 + 3^2 = 54$$

$$SS_y = 4^2 + 6^2 + 5^2 + 7^2 + 8^2 = 190$$

Comments

29

- $2d + 1$ values represent any number of points
 - ▣ d = number of dimensions
- **Average in each dimension j : LS_j/N**
- **Variance in each dimension j :**
 $(SS_j/N) - (LS_j/N)^2$
 - ▣ To get standard deviation take square root
- Can also compute the radius (next slide)

Radius Derivation

30

$$R = \sqrt{\frac{\sum_{i=1}^N \sum_{j=1}^d (x_{ij} - \mu_j)^2}{N}} = \sqrt{\frac{\sum_{j=1}^d \sum_{i=1}^N (x_{ij}^2 - 2x_{ij}\mu_j + \mu_j^2)}{N}}$$

Definition of SS

Definition of LS

$$= \sqrt{\frac{\sum_{j=1}^d (\sum_{i=1}^N x_{ij}^2) - 2\mu_j (\sum_{i=1}^N x_{ij}) + (N\mu_j^2)}{N}}$$

Definition of centroid

$$= \sqrt{\frac{\sum_{j=1}^d SS_j - 2 \frac{LS_j}{N} LS_j + N \left(\frac{LS_j}{N} \right)^2}{N}}$$

Cluster Feature Tree

31

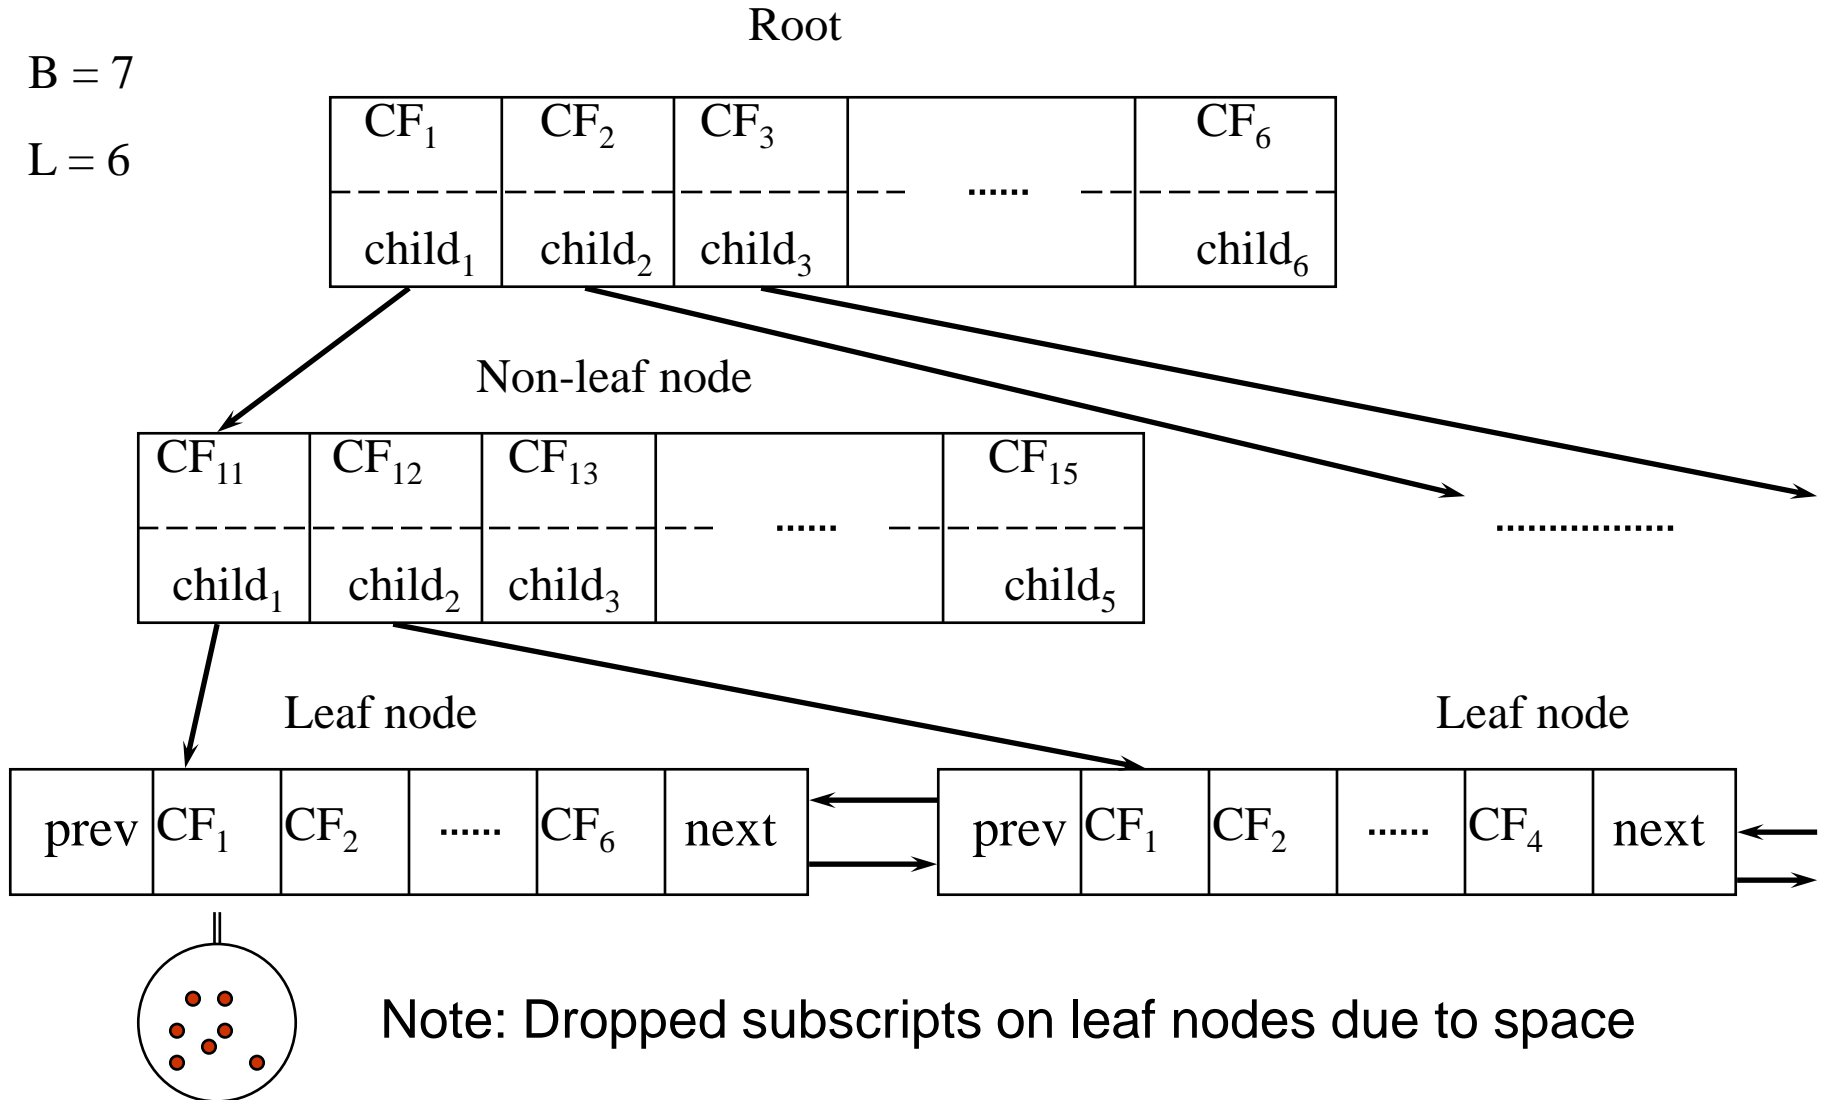
- A CF-tree is a height-balanced tree with two parameters:
 - ▣ Branching factor (non leaf nodes B , leaf nodes, L)
 - ▣ Threshold T
- Each non leaf node has the form $[CF_i, child_i]$
- Each leaf node has CF
 - ▣ Set of CFs
 - ▣ Two pointers: prev and next
- Radius of a subcluster under a leaf node can not exceed the threshold T

CF Tree

32

$B = 7$

$L = 6$



CF-Tree Construction

33

- Scan data set and insert the incoming data instances into the CF tree one by one
- Each instance is inserted into the closest subcluster under a leaf node
- If insertion causes subcluster radius to exceed threshold, then create new subcluster

CF-Tree Construction

34

- ❑ The new subcluster may cause its parent to exceed branching factor
- ❑ If so, split leaf node
 - ▣ Identifying the pair of subclusters with largest inter-cluster distance
 - ▣ Divide by proximity to these two subclusters
- ❑ If this split cause non-leaf node to exceed branching factor, then recursively split
- ❑ If the root node is split, then the height of the CF tree is increased by one


Example: Insert (4,4) into CF Tree

35

Assume
 $B=L=4$
 $R = 2$

(7, (60,60), (726,726))	(2, (1000,1000), (50000,50000))		

(3, (12,12), (50,50))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	
-----------------------------	-------------------------------	-------------------------------	-------------------------------	--



Example: Insert (4,4) into CF Tree

36

(4,4) clearly closest to this CF

(7,
(60,60),
(726,726))

(2,
(1000,1000),
(50000,50000))

Thus sort down to this child

	(3, (12,12), (50,50))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	
--	-----------------------------	-------------------------------	-------------------------------	-------------------------------	--



Example: Insert (4,4) into CF Tree

37

Find CF (4,4) is closest to in this leaf

(7, (60,60), (726,726))	(2, (1000,1000), (50000,50000))		

(3, (12,12), (50,50))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	



Example: Insert (4,4) into CF Tree

38

This one!
Now update the CF!

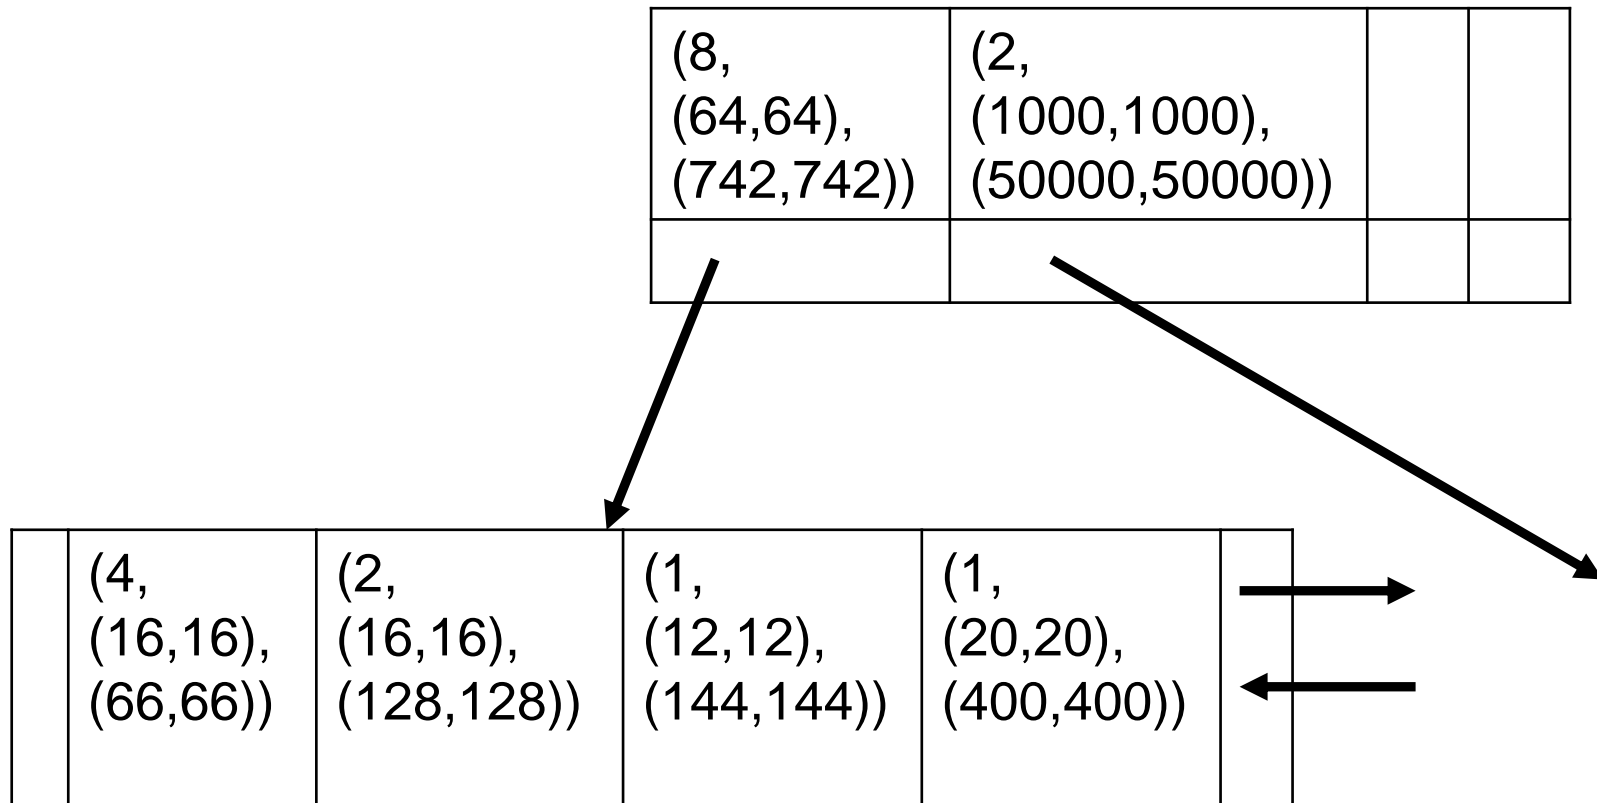
(7, (60,60), (726,726))	(2, (1000,1000), (50000,50000))		

(3, (12,12), (50,50))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	



Example: Result After Inserting (4,4) into CF Tree

39



Now Insert (50,50) into CF Tree

40

Again (50,50)
clearly closest
to this CF

(8, (64,64), (742,742))	(2, (1000,1000), (50000,50000))		

(4, (16,16), (66,66))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	
-----------------------------	-------------------------------	-------------------------------	-------------------------------	--



Now Insert (50,50) into CF Tree

41

Find CF (50,50) is closest to in this leaf

(8, (64,64), (742,742))	(2, (1000,1000), (50000,50000))		

(4, (16,16), (66,66))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	



Now Insert (50,50) into CF Tree

42

Adding (50,50) to
any CF would exceed
Radius

(8, (64,64), (742,742))	(2, (1000,1000), (50000,50000))		

(4, (16,16), (66,66))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	



Now Insert (50,50) into CF Tree

43

Alas, no space to add a new CF, so
Must split


(8, (64,64), (742,742))	(2, (1000,1000), (50000,50000))		

(4, (16,16), (66,66))	(2, (16,16), (128,128))	(1, (12,12), (144,144))	(1, (20,20), (400,400))	




After Inserting (50,50) into CF Tree

44



Traditional Algorithms

45



All-Points Based

d_{\min}, d_{\max}

Centroid Based


$d_{\text{avg}}, d_{\text{mean}}$

What Would BIRCH Do?

46

□ BIRCH assumes:

- ▣ Clusters are normally distributed in each dimension
- ▣ Axes are fixed: Ellipses at an angle are *not OK*




Clustering Using Representatives (CURE)

47

- Cluster definition: Set of representative points
 - ▣ Enables clusters of differing shapes
- Requires an Euclidean space
- Two-pass (hierarchical) clustering approach
 - ▣ Pass 1: Clustering of subset of data to pick “representative” points
 - ▣ Pass 2: Assign all points to clusters

Example: Stanford Salaries

48




Pass 1

49

- Randomly sample of data that fits in memory
- Find initial clusters: Hierarchically cluster the data sample
- For each cluster, pick representative points
 - Select subset of points, as dispersed as possible to represent cluster
 - Move these points towards cluster center (e.g., shrink 20% towards mean)


Example: Initial Clusters

50




Example: Pick Dispersed Points

51



Example: Pick Dispersed Points



52



Pass 2

53


- Scan entire data set
- Assign each example e to “closest” cluster
 - ▣ Standard metric determines closest
 - ▣ Done by finding representative with smallest distance to e



BIRCH vs. CURE Summary


54

BIRCH




Fixed axes, normally
distributed in each
dimension

CURE



Rotated axes



Non-ellipsoid
shape

Summary

55

- Hierarchical clustering models hierarchical structure among the examples
- Typically learned in a bottom-up manner
 - ▣ Linkage and distance are key parameters
 - ▣ Scalability is a key concern
- Advanced algorithms improve efficiency
- Shapes that can be represented are algorithm dependent

Questions?

56

- CURE slides from MMDS.org